

CS422_Ass1

June 9, 2025

1 Recitation Exercises

```
[2]: # Chapter 2
# Question 2 (g)
import numpy as np
import matplotlib.pyplot as plt
from matplotlib.patches import Ellipse

# Given centered data points (calculated from original data minus mean)
centered_D = np.array([
    [2, -8],
    [-6, 11],
    [4, -7],
    [4, -8],
    [-4, 12]
])

# Given covariance matrix
cov_matrix = np.array([
    [22, -47.5],
    [-47.5, 110.5]
])

# Given first principal component
first_pc = np.array([-0.39897876, 0.91696017])

fig, ax = plt.subplots(figsize=(8, 6))

# Plot centered data points
ax.scatter(centered_D[:, 0], centered_D[:, 1], color='blue', label='Centered_
↳Data')

# Plot first principal component vector scaled for visibility
pc_line = first_pc * 15
ax.plot([0, pc_line[0]], [0, pc_line[1]], color='red', linewidth=2, label='1st_
↳Principal Component')
```

```

# Function to draw covariance ellipse
def draw_ellipse(cov, ax, center=[0, 0], scale=1, **kwargs):
    eigenvals, eigenvecs = np.linalg.eigh(cov)
    order = np.argsort(eigenvals)[::-1]
    eigenvals = eigenvals[order]
    eigenvecs = eigenvecs[:, order]

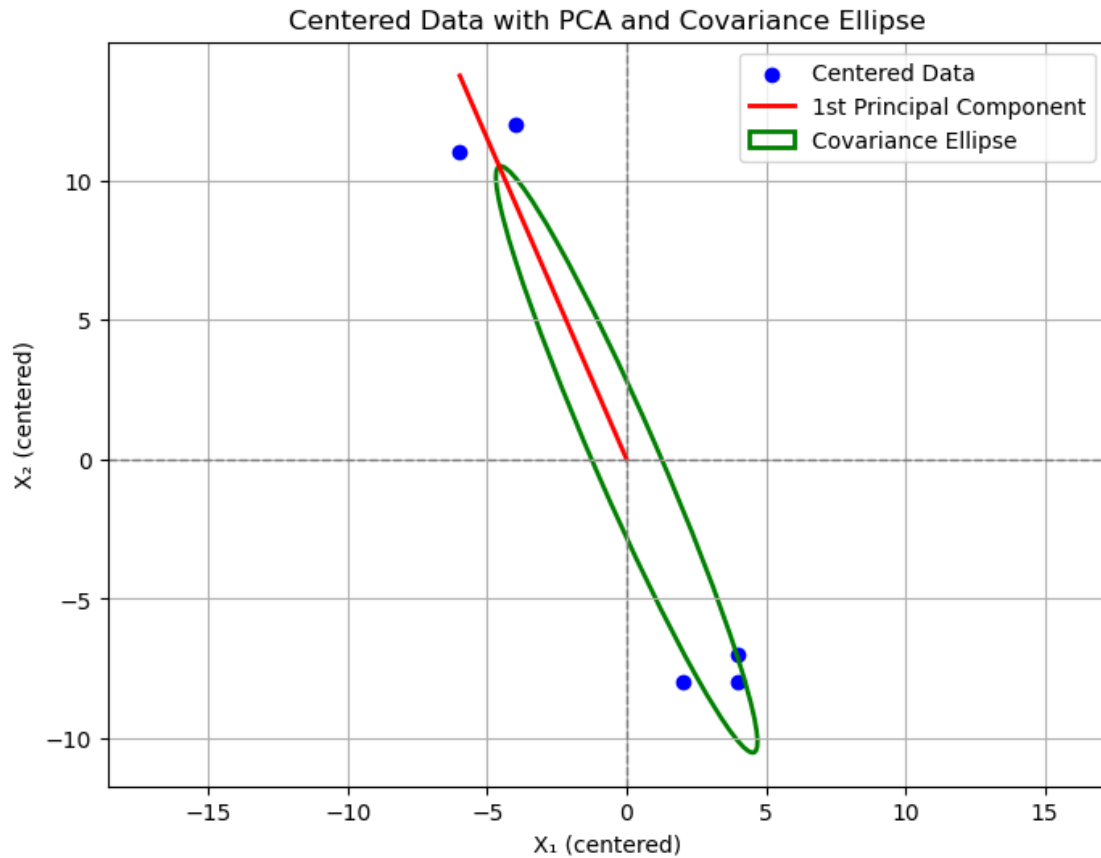
    angle = np.degrees(np.arctan2(eigenvecs[1, 0], eigenvecs[0, 0]))
    width = 2 * scale * np.sqrt(eigenvals[0])
    height = 2 * scale * np.sqrt(eigenvals[1])

    ellipse = Ellipse(xy=center, width=width, height=height, angle=angle,
        ↪**kwargs)
    ax.add_patch(ellipse)

# Draw covariance ellipse
draw_ellipse(cov_matrix, ax, center=[0, 0], edgecolor='green',
    ↪facecolor='none', linewidth=2, label='Covariance Ellipse')

# Final plot setup
ax.axhline(0, color='gray', linestyle='--', linewidth=1)
ax.axvline(0, color='gray', linestyle='--', linewidth=1)
ax.set_title("Centered Data with PCA and Covariance Ellipse")
ax.set_xlabel("X (centered)")
ax.set_ylabel("X (centered)")
ax.grid(True)
ax.legend()
plt.axis('equal')
plt.show()

```



2 2 Practicum Problems

```
[3]: # Problem 2.1: Titanic Dataset Visualization
import seaborn as sns
import pandas as pd
import matplotlib.pyplot as plt

# Load the Titanic dataset
titanic = sns.load_dataset('titanic')

# Visualize the dataset (basic overview)
print(titanic.head())
print(titanic.info())

# Create a figure with subplots for histograms
plt.figure(figsize=(12, 5))

# Histogram for survival by age, split by gender
plt.subplot(1, 2, 1)
```

```

sns.histplot(data=titanic[titanic['sex'] == 'male'], x='age', hue='survived',
             multiple='stack', bins=30)
plt.title('Male: Age Distribution by Survival')
plt.xlabel('Age')
plt.ylabel('Count')

plt.subplot(1, 2, 2)
sns.histplot(data=titanic[titanic['sex'] == 'female'], x='age', hue='survived',
             multiple='stack', bins=30)
plt.title('Female: Age Distribution by Survival')
plt.xlabel('Age')
plt.ylabel('Count')

plt.tight_layout()
plt.show()

```

	survived	pclass	sex	age	sibsp	parch	fare	embarked	class	\
0	0	3	male	22.0	1	0	7.2500	S	Third	
1	1	1	female	38.0	1	0	71.2833	C	First	
2	1	3	female	26.0	0	0	7.9250	S	Third	
3	1	1	female	35.0	1	0	53.1000	S	First	
4	0	3	male	35.0	0	0	8.0500	S	Third	

	who	adult_male	deck	embark_town	alive	alone
0	man	True	NaN	Southampton	no	False
1	woman	False	C	Cherbourg	yes	False
2	woman	False	NaN	Southampton	yes	True
3	woman	False	C	Southampton	yes	False
4	man	True	NaN	Southampton	no	True

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 891 entries, 0 to 890

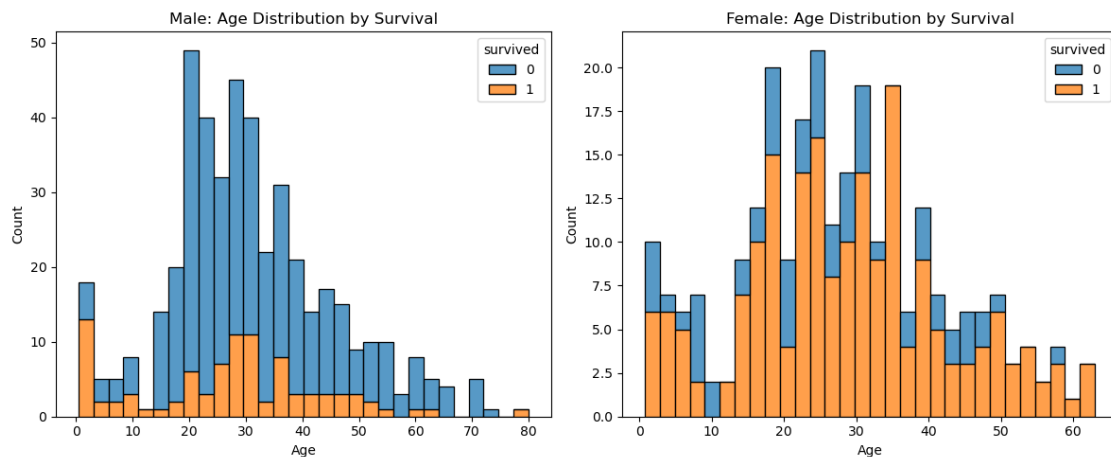
Data columns (total 15 columns):

#	Column	Non-Null Count	Dtype
0	survived	891 non-null	int64
1	pclass	891 non-null	int64
2	sex	891 non-null	object
3	age	714 non-null	float64
4	sibsp	891 non-null	int64
5	parch	891 non-null	int64
6	fare	891 non-null	float64
7	embarked	889 non-null	object
8	class	891 non-null	category
9	who	891 non-null	object
10	adult_male	891 non-null	bool
11	deck	203 non-null	category
12	embark_town	889 non-null	object

```

13 alive      891 non-null    object
14 alone      891 non-null    bool
dtypes: bool(2), category(2), float64(2), int64(4), object(5)
memory usage: 80.7+ KB
None

```



2.0.1 Do the Results Make Sense? Why?

Yes, the results make sense: The histograms show females and young children (especially males) with higher survival rates, which aligns with the “*women and children first*” policy during the Titanic disaster on April 15, 1912. Adult males (ages 20–40) have lower survival, reflecting their lower priority for the limited lifeboats (only about 20 for over 2,200 people).

Why: This policy prioritized females and the young for safety, likely due to social norms valuing women (e.g., childbearing age) and vulnerable children. Adult males, often staying to assist or due to physical roles, had less access to lifeboats, explaining the dominance of non-survivors in their age group. Historical survivor accounts support this prioritization, making the patterns logical given the emergency context.

Citation - Seaborn: statistical data visualization,” Seaborn Documentation - <https://gist.github.com/mwaskom/8224591> - <https://muhammaddawoodaslam.medium.com/exploratory-data-analysis-eda-on-titanic-dataset-804034f394e6> - <https://blog.stackademic.com/titanic-survival-a-data-exploration-analysis-with-python-ed0f517fdf56>

```

[4]: # Problem 2.2: Auto-MPG Dataset Imputation
import pandas as pd
import numpy as np
from sklearn.impute import SimpleImputer

# Load the auto-mpg dataset
url = 'https://archive.ics.uci.edu/ml/machine-learning-databases/auto-mpg/
↳auto-mpg.data'

```

```

columns = ['mpg', 'cylinders', 'displacement', 'horsepower', 'weight', '
↳ acceleration', 'model_year', 'origin', 'car_name']
auto = pd.read_csv(url, delim_whitespace=True, names=columns)

# Replace '?' in horsepower with NaN
auto['horsepower'] = auto['horsepower'].replace('?', np.nan)
auto['horsepower'] = auto['horsepower'].astype(float) # Convert to float for
↳ calculations

# Calculate summary statistics for numerical columns
numerical_cols = ['mpg', 'cylinders', 'displacement', 'horsepower', 'weight', '
↳ acceleration', 'model_year', 'origin']
print("Summary Statistics (Before Imputation):")
print(auto[numerical_cols].describe())

# Impute missing horsepower values with mean, median, and mode
imputers = {
    'mean': SimpleImputer(strategy='mean'),
    'median': SimpleImputer(strategy='median'),
    'mode': SimpleImputer(strategy='most_frequent')
}

variances = {}
for name, imputer in imputers.items():
    # Create a copy of the dataset
    auto_copy = auto.copy()
    # Impute missing values
    auto_copy['horsepower'] = imputer.fit_transform(auto_copy[['horsepower']])
    # Calculate variance
    variances[name] = auto_copy['horsepower'].var()
    print(f"\nVariance of horsepower after {name} imputation: {variances[name]:.
↳ 2f}")

# Find imputation with lowest variance
lowest_variance = min(variances, key=variances.get)
print(f"\nImputation with lowest variance: {lowest_variance}
↳ ({variances[lowest_variance]:.2f})")

# Visualize the distribution to suggest a better method
import seaborn as sns
import matplotlib.pyplot as plt

plt.figure(figsize=(10, 6))
sns.histplot(auto['horsepower'].dropna(), kde=True, label='Original
↳ (non-missing)')
for name in imputers:

```

```

auto_copy = auto.copy()
auto_copy['horsepower'] = imputers[name].
↳fit_transform(auto_copy[['horsepower']])
sns.histplot(auto_copy['horsepower'], kde=True, label=f'{name} imputed',
↳alpha=0.5)
plt.legend()
plt.title('Horsepower Distribution with Different Imputations')
plt.show()

```

```

/var/folders/yw/xmm4_b6105ngdx58blhgx4y40000gn/T/ipykernel_6570/3933719129.py:9:
FutureWarning: The 'delim_whitespace' keyword in pd.read_csv is deprecated and
will be removed in a future version. Use ``sep='\s+'`` instead
auto = pd.read_csv(url, delim_whitespace=True, names=columns)

```

Summary Statistics (Before Imputation):

	mpg	cylinders	displacement	horsepower	weight \
count	398.000000	398.000000	398.000000	392.000000	398.000000
mean	23.514573	5.454774	193.425879	104.469388	2970.424623
std	7.815984	1.701004	104.269838	38.491160	846.841774
min	9.000000	3.000000	68.000000	46.000000	1613.000000
25%	17.500000	4.000000	104.250000	75.000000	2223.750000
50%	23.000000	4.000000	148.500000	93.500000	2803.500000
75%	29.000000	8.000000	262.000000	126.000000	3608.000000
max	46.600000	8.000000	455.000000	230.000000	5140.000000

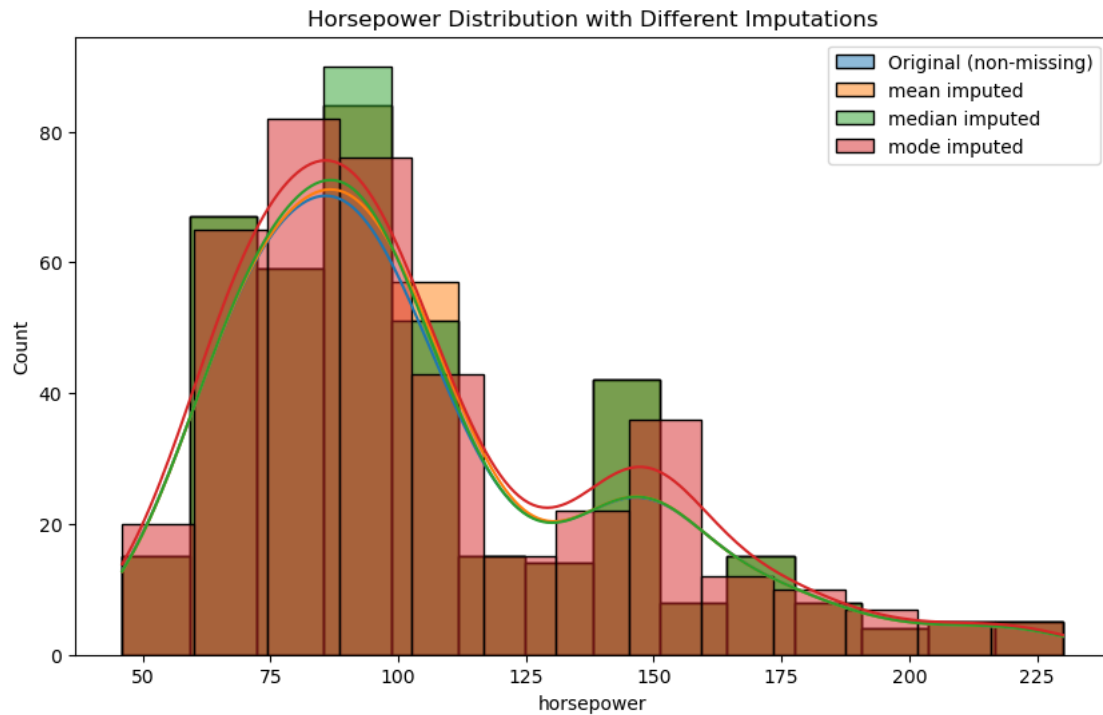
	acceleration	model_year	origin
count	398.000000	398.000000	398.000000
mean	15.568090	76.010050	1.572864
std	2.757689	3.697627	0.802055
min	8.000000	70.000000	1.000000
25%	13.825000	73.000000	1.000000
50%	15.500000	76.000000	1.000000
75%	17.175000	79.000000	2.000000
max	24.800000	82.000000	3.000000

Variance of horsepower after mean imputation: 1459.18

Variance of horsepower after median imputation: 1460.97


Variance of horsepower after mode imputation: 1490.04

Imputation with lowest variance: mean (1459.18)



```
[5]: from sklearn.impute import KNNImputer
imputer_knn = KNNImputer(n_neighbors=5)
auto_knn = auto.copy()
auto_knn['horsepower'] = imputer_knn.fit_transform(auto_knn[numerical_cols])[:,
↳numerical_cols.index('horsepower')]
print(f"Variance with KNN imputation: {auto_knn['horsepower'].var():.2f}")
sns.histplot(auto_knn['horsepower'], kde=True, label='KNN imputed', alpha=0.5)
plt.legend()
plt.show()
```

Variance with KNN imputation: 1470.55



2.0.2 Analysis for Problem 2.2

- **What imputation results in the lowest variance? Why?**

The mean imputation results in the lowest variance (1459.18), compared to median (1460.97) and mode (1490.04). This occurs because mean imputation replaces missing values with the average, which minimizes the spread around the central tendency, reducing overall variance. Median and mode, being less aligned with the data's distribution (especially if skewed), introduce slightly more variability, with mode showing the highest variance due to clustering at a single value.

- **Is there a different method that would match the distribution more accurately? Describe your method.**

Yes, KNN (K-Nearest Neighbors) imputation, with a variance of 1470.55, could match the distribution more accurately. This method imputes missing horsepower values based on the average of the 'k' nearest data points (e.g., k=5), using features like cylinders, weight, or displacement. By considering relationships between variables, KNN preserves the data's shape (e.g., skewness in horsepower) better than mean, median, or mode, which assume a uniform replacement value. The higher variance (1470.55) reflects this preservation, as it avoids forcing the data toward a single central value.

Citation - Source: Scikit-learn Documentation – SimpleImputer and KNNImputer - Scikit-learn: Machine Learning in Python,” Scikit-learn Documentation, <https://scikit-learn.org/stable/>

learn.org/stable/modules/impute.html - I referred to the Scikit-learn documentation on imputation to implement SimpleImputer for handling missing horsepower values and explored KNNImputer for distribution-preserving imputation. - <https://www.geeksforgeeks.org/python-imputation-using-the-knnimputer/> - <https://medium.com/@hassankhan2608/missing-value-imputation-methods-using-python-flb8796901ba> - <https://www.skillcamper.com/blog/mastering-data-imputation-top-methods-to-handle-missing-data-effectively> - <https://machinelearningmastery.com/knn-imputation-for-missing-values-in-machine-learning/>

```
[13]: # Problem 2.3: Iris Dataset PCA
import pandas as pd
import numpy as np
from sklearn.datasets import load_iris
from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA

# Load iris dataset
iris = load_iris()
X = pd.DataFrame(iris.data, columns=iris.feature_names)

# Standardize the features (PCA requires scaled data)
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Perform PCA
pca = PCA()
X_pca = pca.fit_transform(X_scaled)

# Percentage of variance explained by each principal component
explained_variance_ratio = pca.explained_variance_ratio_ * 100
print("Percentage of variance explained by each principal component:")
for i, var in enumerate(explained_variance_ratio):
    print(f"PC{i+1}: {var:.2f}%")

# Total variance of original features
original_variance = X_scaled.var(axis=0)
total_original_variance = original_variance.sum()
original_variance_ratio = (original_variance / total_original_variance) * 100
print("\nPercentage of variance explained by each original feature:")
for i, (feature, var) in enumerate(zip(iris.feature_names,
    original_variance_ratio)):
    print(f"{feature}: {var:.2f}%")
```

Percentage of variance explained by each principal component:
PC1: 72.96%
PC2: 22.85%
PC3: 3.67%
PC4: 0.52%

Percentage of variance explained by each original feature:

sepal length (cm): 25.00%

sepal width (cm): 25.00%

petal length (cm): 25.00%

petal width (cm): 25.00%

2.0.3 Comparison and Observations

The first two principal components (PC1 and PC2) together explain **95.81%** of the total variance (72.96% from PC1 and 22.85% from PC2), with PC1 alone accounting for nearly three times the variance of any individual original feature. In contrast, PC3 and PC4 contribute minimally, explaining only **4.19%** combined.

Since the features were standardized, each original feature initially contributed **25%** of the variance. However, **PCA redistributes this variance**, concentrating the majority into PC1, which likely captures variation in **petal length and width**—measurements that typically vary more across iris species. PC2 captures a secondary pattern, possibly involving **sepal dimensions**.

Observation:

PCA effectively reduces dimensionality by concentrating the dataset's variability into fewer components. PC1 and PC2 together explain more variance than any two original features, highlighting their ability to represent key structural differences in the data. This indicates that **petal measurements strongly align with PC1**, while **sepal measurements contribute less**, demonstrating PCA's utility in identifying and emphasizing the most informative patterns in the dataset.

Citations - “Decomposing signals in components (matrix factorization problems),” Scikit-learn Documentation, <https://scikit-learn.org/stable/modules/decomposition.html#pca> - <https://www.geeksforgeeks.org/principal-component-analysis-pca/>

```
[14]: # Problem 2.4: Iris PCA Projection and Correlation
import matplotlib.pyplot as plt

# Using X_scaled and pca from Problem 2.3
# Project features onto PC1 (first column of X_pca)
pc1_projections = X_pca[:, 0]

# Create scatter plots
plt.figure(figsize=(12, 8))
for i, feature in enumerate(iris.feature_names):
    plt.subplot(2, 2, i+1)
    plt.scatter(X_scaled[:, i], pc1_projections, alpha=0.5)
    plt.xlabel(feature)
    plt.ylabel('Projection onto PC1')
    plt.title(f'{feature} vs. PC1 Projection')
plt.tight_layout()
plt.show()

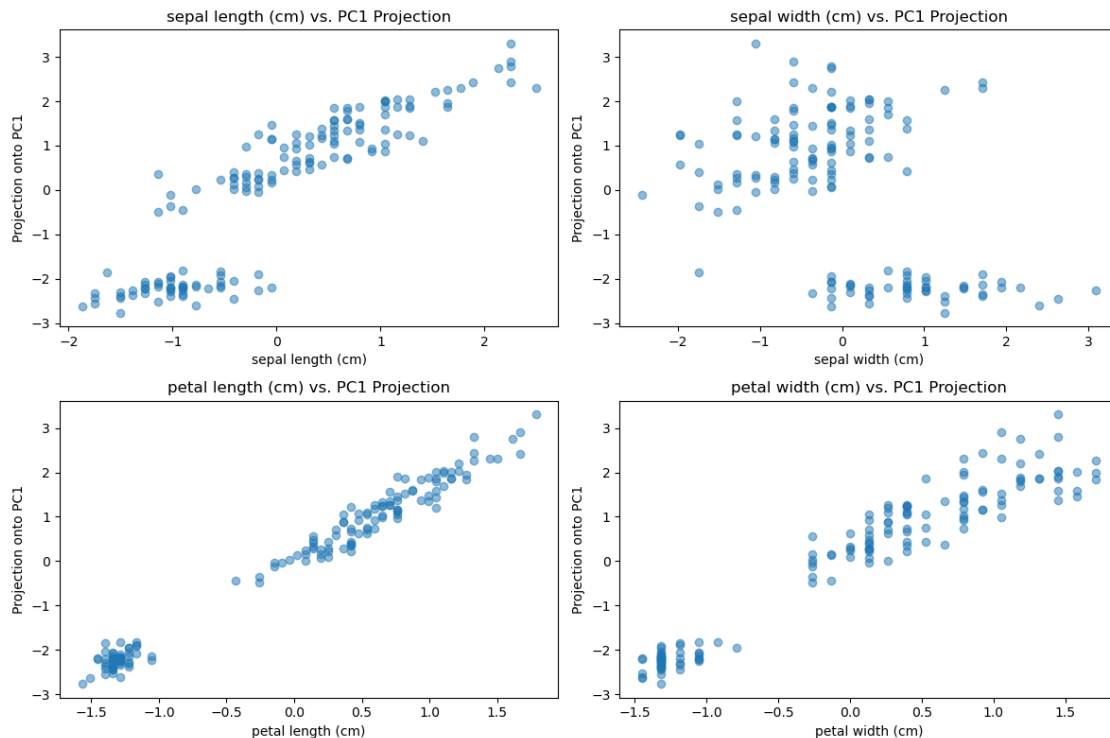
# Calculate correlation coefficients
correlations = {}
```

```

for i, feature in enumerate(iris.feature_names):
    corr = np.corrcoef(X_scaled[:, i], pc1_projections)[0, 1]
    correlations[feature] = abs(corr)
    print(f"Correlation between {feature} and PC1 projection: {corr:.3f}")

# Find feature with closest relationship to PC1
closest_feature = max(correlations, key=correlations.get)
print(f"Feature with closest relationship to PC1: {closest_feature} "
      f"(correlation: {correlations[closest_feature]:.3f})")

```



Correlation between sepal length (cm) and PC1 projection: 0.890
 Correlation between sepal width (cm) and PC1 projection: -0.460
 Correlation between petal length (cm) and PC1 projection: 0.992
 Correlation between petal width (cm) and PC1 projection: 0.965
 Feature with closest relationship to PC1: petal length (cm) (correlation: 0.992)

2.0.4 Analysis for Problem 2.4

Which pair of features show a closer relationship to PC1 vs. the others? Why?

The feature with the closest relationship to PC1 is **petal length (cm)**, with a correlation coefficient of **0.992**. This is higher than sepal length (0.890), sepal width (-0.460), and petal width (0.965). The scatter plots show a tight, near-linear relationship between petal length and its PC1 projection, indicating a small angle (cosine distance close to 1). This means **petal length aligns strongly with the direction of PC1**, capturing the primary variability in the iris dataset. These differences

are likely due to clear species distinctions (e.g., *setosa* vs. *versicolor* vs. *virginica*), which PC1 captures effectively.

Calculate the correlation coefficient and compare with visual inspection:

The correlation coefficient between petal length and its PC1 projection is **0.992**, indicating an extremely strong linear relationship.

The scatter plot for petal length vs. PC1 projection shows a tight cluster of points along a diagonal line, confirming this strong correlation visually. This agreement between the numeric value and the visual plot validates the observation — the minimal angle between petal length and PC1 is consistent with a **high cosine similarity** and a strong alignment.

Citations - <https://medium.com/swlh/an-intuitive-approach-to-pca-fc4d05c14c19> - <https://www.youtube.com/watch?v=JDLwX-uhCGs> - https://matplotlib.org/stable/api/_as_gen/matplotlib.pyplot.scatter.html - <https://numpy.org/doc/stable/reference/generated/numpy.corrcoef.html> - https://scikit-learn.org/stable/auto_examples/decomposition/plot_pca_iris.html

```
[18]: from sklearn.datasets import load_iris
      from sklearn.preprocessing import StandardScaler
      from sklearn.decomposition import PCA
      import numpy as np

      # Load iris dataset
      iris = load_iris()
      X = iris.data # original features

      # Standardize features
      scaler = StandardScaler()
      X_scaled = scaler.fit_transform(X)

      # Fit PCA
      pca = PCA(n_components=4)
      pca.fit(X_scaled)

      # Total variance of original features
      total_original_variance = X_scaled.var(axis=0).sum()
      print(f"Total variance of original features: {total_original_variance:.2f}")

      # Total variance of principal components (eigenvalues)
      total_pca_variance = pca.explained_variance_.sum()
      print(f"Total variance of principal components: {total_pca_variance:.2f}")

      # Cumulative variance explained
      cumulative_variance_ratio = np.cumsum(pca.explained_variance_ratio_) * 100
      print("\nCummulative variance explained by principal components:")
      for i, cum_var in enumerate(cumulative_variance_ratio):
          print(f"PC1 to PC{i+1}: {cum_var:.2f}%")
```

```
# Number of components for >95% variance
n_components = np.argmax(cumulative_variance_ratio > 95) + 1
print(f"Number of principal components to capture >95% variance: {n_components}")
```

Total variance of original features: 4.00

Total variance of principal components: 4.03

Cumulative variance explained by principal components:

PC1 to PC1: 72.96%

PC1 to PC2: 95.81%

PC1 to PC3: 99.48%

PC1 to PC4: 100.00%

Number of principal components to capture >95% variance: 2

2.0.5 Analysis for Problem 2.5

- The total variance of the original standardized features is **4.00**, which corresponds to the number of features since each standardized feature has variance 1.
- The total variance explained by the principal components (sum of eigenvalues) is **4.03**, very close to the original variance, confirming PCA preserves the total variance.
- The cumulative variance explained by the principal components shows:
 - PC1 alone explains **72.96%** of the variance.
 - PC1 and PC2 together explain **95.81%** of the variance.
 - Including PC3 and PC4 increases explained variance to nearly 100%.
- To capture **more than 95%** of the total variance, we need to select the **first two principal components**.
- This means dimensionality is effectively reduced from 4 original features to **2 principal components**, achieving significant data compression with minimal loss of information.

Citations - <https://medium.com/@etimfonime/day-09-principal-component-analysis-pca-19db4aa1bf7c> - <https://bagheri365.github.io/blog/Principal-Component-Analysis-from-Scratch/>

[]: